

# Zero Downtime Migrations in Django

---

*Ryan Scott*

# Introduction

- Software Engineer @ Percipient Networks
- Several years using Python and Django for fun
- 7 months using Python and Django professionally
- ~1 year using Python and Django professionally

# Topics

- What does zero downtime mean?
- What is a zero downtime deployment?
- What is a migration?
- How does Django query for data?
- How do you migrate without downtime?



# What does zero downtime mean?

- End-users can continue using the website without noticing anything while you deploy.
- All services remain functioning.
- No downtime when you make changes!

# Zero downtime deployments

- Without zero downtime deployments, zero downtime migrations are useless.
- Deploying without downtime requires multiple Django servers (load balanced).
- Rolling deploy: Deploy to one web server at a time
- Blue/green deploy: Stand up a duplicate set of web servers, swap them out after the deploy has finished

# What's a migration?

- Migrations allow Django to update the database schema from changes you've made to models.
- Migrations maintain a history of your database state.
- Migrations produce consistent results so that you can mirror development and production database state.

# Zero downtime migrations

- How do you make database changes without breaking Django?
- How do you keep your Django app running while changing the database?



## Example model

- A Customer represents a company with a name and a plan.

```
class Customer(models.Model):  
    company_name = models.CharField(max_length=32)  
    plan = models.CharField(max_length=32) # paid, free, etc.
```

# How does Django query for data?

- Django query:

```
Customer.objects.all()
```

- Translates to this SQL query:

```
SELECT company_name, plan FROM customer;
```

- NOT:

```
SELECT * FROM customer;
```

## Leveraging Django's querying method

- Django only queries for fields that it knows about.
- Fields can exist in the database without being defined in Django.
- Django will raise an exception if a field is defined, but isn't in the database.

# Migration Strategies

- Adding a field
- Removing a field

# Adding a field

- Add a new field for the company's address. The field *can* be blank.

```
class Customer(models.Model):  
    company_name = models.CharField(max_length=32)  
    plan = models.CharField(max_length=32) # paid, free, etc.  
    company_address = models.CharField(blank=True, max_length=100)
```

# Adding a field

Steps to deploy the new field:

1. Migrate the database so that the new field exists.
2. Deploy your code to Django.

This ensures that the field exists in the database *before* Django starts using it.

# Removing a field

- Remove the plan field because we decided we don't need it

```
class Customer(models.Model):  
    company_name = models.CharField(max_length=32)  
    # Removed plan field  
    company_address = models.CharField(blank=True, max_length=100)
```

# Removing a field

Steps to deploy the deleted field:

1. Deploy your code to Django.
2. Migrate the database so that the old field is deleted.

This ensures that Django stops using the field *before* it is deleted from the database.



# Complex migration strategies

- RunPython can be used in Django migrations for migrating data (backfilling).
- Exceptions in RunPython functions can leave your migration in a weird state.
- Management commands often work better for data migrations because they are easier to test and can be run repeatedly.

# Complex migration strategies

- More complex migrations can be designed by building on these basic steps.
- *Remember:* Understanding how Django queries for data will help you determine the order for migrating and deploying.
- Testing complex migrations:  
[github.com/plumdog/  
django\\_migration\\_testcase](https://github.com/plumdog/django_migration_testcase)

Thank you

[github.com/percipient/talks](https://github.com/percipient/talks)

[ryan@strongarm.io](mailto:ryan@strongarm.io)