

TWISTED

CONCEPTS & PATTERNS

INTRODUCTION

PATRICK CLOKE

LEAD ENGINEER @
PERCIPIENT NETWORKS

&

MOZILLIAN
[GITHUB.COM/CLOKEP](https://github.com/clokep)





STEPHEN DICATO

**CO-FOUNDER & VP,
ENGINEERING**

@ PERCIPIENT NETWORKS

GITHUB.COM/DICATO

IF YOU TWITTER
**@CLOKEP &
@STEPHENDICATO**



STRONGARM.IO
automate your incident response

questions
WELCOME



[GITHUB.COM/PERCIPIENT/TALKS](https://github.com/perceptient/talks)

expectations

LEARN THESE 6 CONCEPTS

1. WHAT IS ASYNC PROGRAMMING

2. WHAT IS TWISTED

3. WHEN/WHY TO USE TWISTED

4. EVENT LOOP (REACTOR)

5. DEFERREDS

6. PROTOCOLS (AND MORE)

**BONUS: USE TWISTED TO
BUILD SYSTEMS & SERVICES**

COMPOSABLE & SCALABLE SYSTEMS

**I.E., SERVICE ORIENTED
ARCHITECTURE**


I.E., MICROSERVICES


CONTEXT

AN EXAMPLE APP

chat server **EXAMPLE (WITH ADMIN DASHBOARD)**


- ▶ **Clients use netcat**
- ▶ **Messages are broadcast to all users**
 - ▶ **Clients are sent a banner on login**
- ▶ **Banner is configurable via an admin webpage**






1. WHAT IS ASYNC PROGRAMMING

SYNCHRONOUS




<http://krondo.com/>

THREADED



<http://krondo.com/>


MULTIPROCESSING



multiprocessing uses same model


<http://krondo.com/>

ASYNC



<http://krondo.com/>

BLOCKING



<http://krondo.com/>

2. WHAT IS TWISTED

TWISTED:
AN EVENT-DRIVEN,
asynchronous
NETWORKING ENGINE

EVENT-DRIVEN & ASYNCRHONOUS

- ▶ **user code is triggered when something it cares about happens**
 - ▶ **examples: new connection, data is available, HTTP request received, etc.**
- ▶ **non-blocking while waiting for I/O (e.g. bytes on a socket)**

NETWORKING ENGINE

- ▶ supports many clients & servers out of the box
e.g. web (HTTP), chat (IRC, XMPP), mail (SMTP/IMAP/POP3)
- ▶ easily develop fully custom protocols
- ▶ easily customize behavior of built in protocols

3. WHEN/WHY TO USE TWISTED

1. **high-level networking APIs**
2. **large protocol library**
3. **scalable due to asynchronous design**
4. **mature, with an emphasis on quality**

TWISTED WILL NOT

- ▶ magically make code *non-blocking*
i.e. help with CPU-bound tasks[†]
- ▶ be the simplest library to make a simple HTTP request[‡]

4. EVENT LOOP (REACTOR)

EVENT LOOP APIS:

**NETWORKING, THREADING, EVENT
DISPATCHING, TIMING, ETC.**

**REACTOR DEPENDS ON
PLATFORM AND USAGE
(REACTOR IS A *global singleton*)**

EXAMPLE (FROM runner.py)

```
from twisted.internet import reactor, endpoints
from twisted.web import server
```

```
from chat import NetCatChatFactory
```

```
# Create an instance of the factory.
factory = NetCatChatFactory()
```

```
# Listen on TCP port 1400 for chat.
endpoints.serverFromString(reactor, "tcp:1400").listen(factory)
```

```
# Start listening for connections (and run the event-loop).
reactor.run()
```


```
# Note that any code after this point will *not* be executed. reactor.run enters
# an infinite loop until shutdown.
```

5. DEFERREDS

A *deferred* IS A PROMISE
THAT A FUNCTION WILL EVENTUALLY HAVE A
result

IN OTHER WORDS:
DEFERREDS ARE A PLACEHOLDER FOR A
FUTURE RESULT

DEFERREDS MANAGE A CALLBACK CHAIN



EXAMPLE (FROM deferred_ex.py)

```
import json
from twisted.internet import reactor
from twisted.python import log
from twisted.web import client


def gotOrgs(data):
    print("Organization request done!")
    # Parse the JSON payload. TODO Error checking.
    orgs = json.loads(data)
    # Find the names of the organizations and print them.
    org_names = sorted([org['login'] for org in orgs])
    print('\n'.join(org_names))

def shutdown(ignored):
    print("Shutting down!")
    reactor.stop() # No matter what happens, shutdown the eventloop.

# The Deferred.
d = client.getPage('https://api.github.com/users/clokep/orgs')
d.addCallback(gotOrgs) # The callback for a successful request.
d.addErrback(log.err) # Before shutdown, log any errors.
d.addBoth(shutdown) # The callback/errback.

reactor.run() # Start the eventloop.
```

EXAMPLE DEFERRED FLOW (FROM deferred_ex.py)



DEMO

ADVANCED EXAMPLE, PART 1. (FROM deferred_ex_2.py)

```
import json
from twisted.internet import defer, reactor
from twisted.python import log
from twisted.web import client

def gotRepos(data, org):
    pass

def gotOrgs(data):
    pass

def printOrgs(org_list):
    print("Outputting repos")
    for success, (org, repos) in org_list:
        print('\t%s: %s' % (org, ', '.join(repos) if repos else '(none)'))

def shutdown(ignored):
    print("Shutting down!")
    reactor.stop() # No matter what happens, shutdown the eventloop.

# The Deferred.
d = client.getPage('https://api.github.com/users/clokep/orgs')
d.addCallback(gotOrgs) # The callback for a successful request.
d.addCallback(printOrgs)
d.addErrback(log.err) # Before shutdown, log any errors.
d.addBoth(shutdown) # The callback/errback.

reactor.run() # Start the eventloop.
```

ADVANCED EXAMPLE, PART 2. (FROM deferred_ex_2.py)

```
def gotRepos(data, org):
    """Got the repos for an org, return a tuple of (org name, repos)."""
    print("Got repo information for %s" % org)

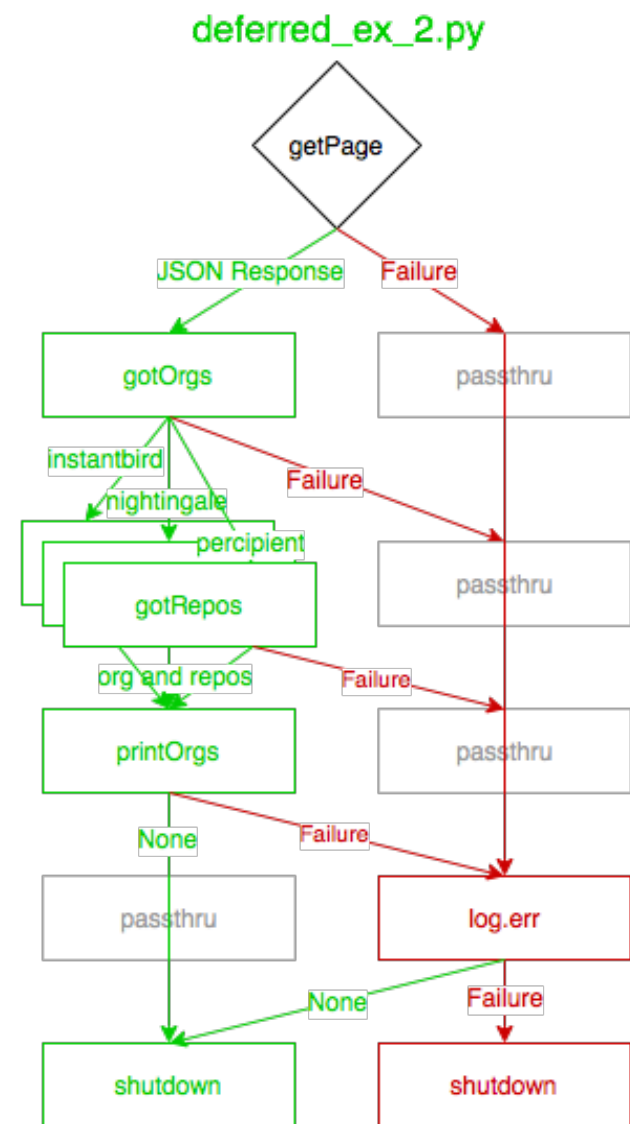
    # Parse the JSON payload. TODO Error checking.
    repos = json.loads(data)
    if not repos: # no repos, return early
        return (org, [])
    # The names of the repos in alphabetical order.
    names = sorted([repo['name'] for repo in repos])
    return (org, names)

def gotOrgs(data):
    print("Organization request done!")
    # Parse the JSON payload. TODO Error checking.
    orgs = json.loads(data)
    # The names of the organizations in alphabetical order.
    org_names = sorted([bytes(org['login']) for org in orgs])

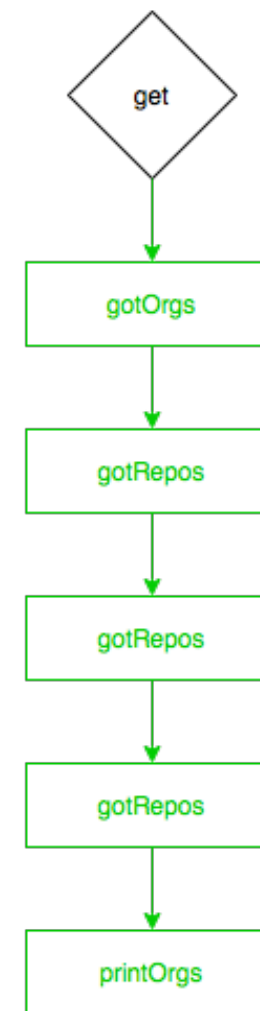
    # Now request the repos under each org.
    ds = []
    for org in org_names:
        print("\t%s" % org) # print out the org
        d = client.getPage('https://api.github.com/orgs/%s/repos' % org)
        d.addCallback(gotRepos, org) # pass the org name to the callback.
        ds.append(d)

    # Returning a Deferred causes the next callback to wait.
    return defer.DeferredList(ds)
```

ADVANCED EXAMPLE DEFERRED FLOW (FROM deferred_ex.py)



deferred_ex_2_req.py




DEMO

DEFERREDS FINE PRINT

```
d = defer.Deferred()
d.addCallback(callback_1)
d.addErrback(errback_1)
d.addCallbacks(callback_2, errback_2)
d.addBoth(cleanup)
```

1. Each added callback/errback appends a new *level* to the chain
2. Errors in callbacks propagate to the next errback
3. Propagation from the errback chain to the callback chain is possible



6. PROTOCOLS (AND MORE)

Protocols: EVENT HANDLERS FOR A CONNECTION

- ▶ Each new connection gets a new Protocol *instance*
 - ▶ Basic events include:
connection opened/closed, data available
- ▶ Transforms wire protocol into higher level events
(e.g. *Line* received or *HTTP request* finished)

EXAMPLE (FROM chat.py)

```
from twisted.internet import protocol

class NetCatChatProtocol(protocol.Protocol):
    # An instance of a Protocol exists for each established connection.

    def connectionMade(self):
        # Called when the protocol is instantiated and the connection is ready.

    def dataReceived(self, data):
        # New data (bytes) are available for consuming.

    def connectionLost(self, reason):
        # The connection is about to be terminated.

# Has a transport property for interacting with the connection.

# Has a factory property for interacting with the factory that build this.
```

ProtocolFactory

KEEPS *state* ACROSS Protocols

BUILDS Protocol INSTANCES

EXAMPLE (FROM chat.py)

```
from twisted.internet import protocol

class NetCatChatProtocol(protocol.Protocol):
    # See above.

class NetCatChatFactory(protocol.Factory):
    # By defining `protocol`, the default implementation of
    # `Factory.buildProtocol` will work fine!
    protocol = NetCatChatProtocol

    # State and other variables would be stored on the factory.
```

Transport

A WAY TO *send data*

- ▶ **Write *bytes* to a connection/datagram**
 - ▶ **Close a connection**
 - ▶ **Query local/remote addresses**
- ▶ **Do not assume *when* data will be sent**

EXAMPLE (FROM chat.py)


```
from twisted.internet import protocol


class NetCatChatProtocol(protocol.Protocol):
    # An instance of a Protocol exists for each established connection.

    def connectionMade(self):
        # The connection has been established, perform greetings here.
        self.transport.write(self.factory.banner)
```

DEMO


**INTEGRATE TWISTED WITH
OTHER SERVICES**







ADMIN CONSOLE

- 1. current user count**
- 2. set banner**







BUILD THE CHAT SERVER API

- 1. GET /users/ --> returns user count**
- 2. GET /banner/ --> returns the current banner**
- 3. POST /banner/ --> sets the current banner**


```
from twisted.web import resource

class ApiResource(resource.Resource):
    def __init__(self, chat_factory, *args, **kwargs):
        # This needs a reference to the NetCatChat factory object.
        self.chat_factory = chat_factory

        resource.Resource.__init__(self, *args, **kwargs)
```

USER API ENDPOINT

```
class Users(ApiResource):  
    isLeaf = True  
  
    def render_GET(self, request):  
        # The user count from server.  
        user_count = len(self.chat_factory.clients)  
        result = {'users': user_count}  
  
        return json.dumps(result, indent=4, separators=(',', ': ')) + "\n"
```

BANNER API ENDPOINT

```
class Banner(ApiResource):  
    isLeaf = True  
  
    def render_GET(self, request):  
        # Get the banner.  
        result = {'banner': self.chat_factory.banner}  
        return json.dumps(result, indent=4, separators=(',', ': ')) + "\n"
```

```
def _set_banner(self, banner):
    # ... error handling ;- )
    self.chat_factory.banner = bytes(banner) # unicode to bytes

def render_POST(self, request):
    # Set the banner.
    status = "ERROR"
    try:
        content = request.content.read()
        data = json.loads(content)['banner']

        # Make this a Deferred so the function can immediately return.
        d = task.deferLater(reactor, 0.1, self._set_banner, data)
        d.addErrback(log.err)

        status = "SUCCESS"
    except Exception:
        # ... error handling ;- )

    return json.dumps({'status': status})
```

UPDATE REACTOR TO EXPOSE API

```
from twisted.internet import reactor, endpoints
from twisted.web import server

from chat import NetCatChatFactory
from api import Root

# Create an instance of the factories.
factory = NetCatChatFactory()
site = server.Site(Root(factory))

# Listen on TCP port 1400 for chat and port 8080 for the API.
endpoints.serverFromString(reactor, "tcp:1400").listen(factory)
endpoints.serverFromString(reactor, "tcp:8080").listen(site)

# Start listening for connections (and run the event-loop).
reactor.run()
# Note that any code after this point will *not* be executed. reactor.run enters
# an infinite loop until shutdown.
```

**WHAT ABOUT THE
CLIENT CODE?**

API CLIENT CODE / USER COUNT

```
def get_user_count():  
    url = API_URL + '/users/'  
    try:  
        r = requests.get(url)  
    except Exception as e: # Gotta catch 'em all!  
        return None  
  
    result = r.json()  
    return result['users']
```

API CLIENT CODE / GET BANNER

```
def get_banner():  
    url = API_URL + '/banner/'  
    try:  
        r = requests.get(url)  
    except Exception as e: # Gotta catch 'em all!  
        return None  
  
    result = r.json()  
    return result['banner']
```


API CLIENT CODE / SET BANNER

```
def set_banner(banner):  
    url = API_URL + '/banner/'  
    data = json.dumps({'banner': banner})  
    try:  
        r = requests.post(url, data=data)  
    except Exception as e: # Gotta catch 'em all!  
        return None
```

NOW SERVE IT USING

FLASK

```
app = Flask(__name__)
API_URL = 'http://127.0.0.1:8080'

@app.route("/")
def splash():
    user_count = get_user_count()
    banner = get_banner()
    return render_template('index.html', user_count=user_count, banner=banner)

@app.route('/set_banner/', methods=['POST'])
def banner():
    banner = request.values.get('banner', "")
    set_banner(banner)
    return redirect(url_for('splash'))

if __name__ == "__main__":
    app.run()
```

TANGENT: `twistd`

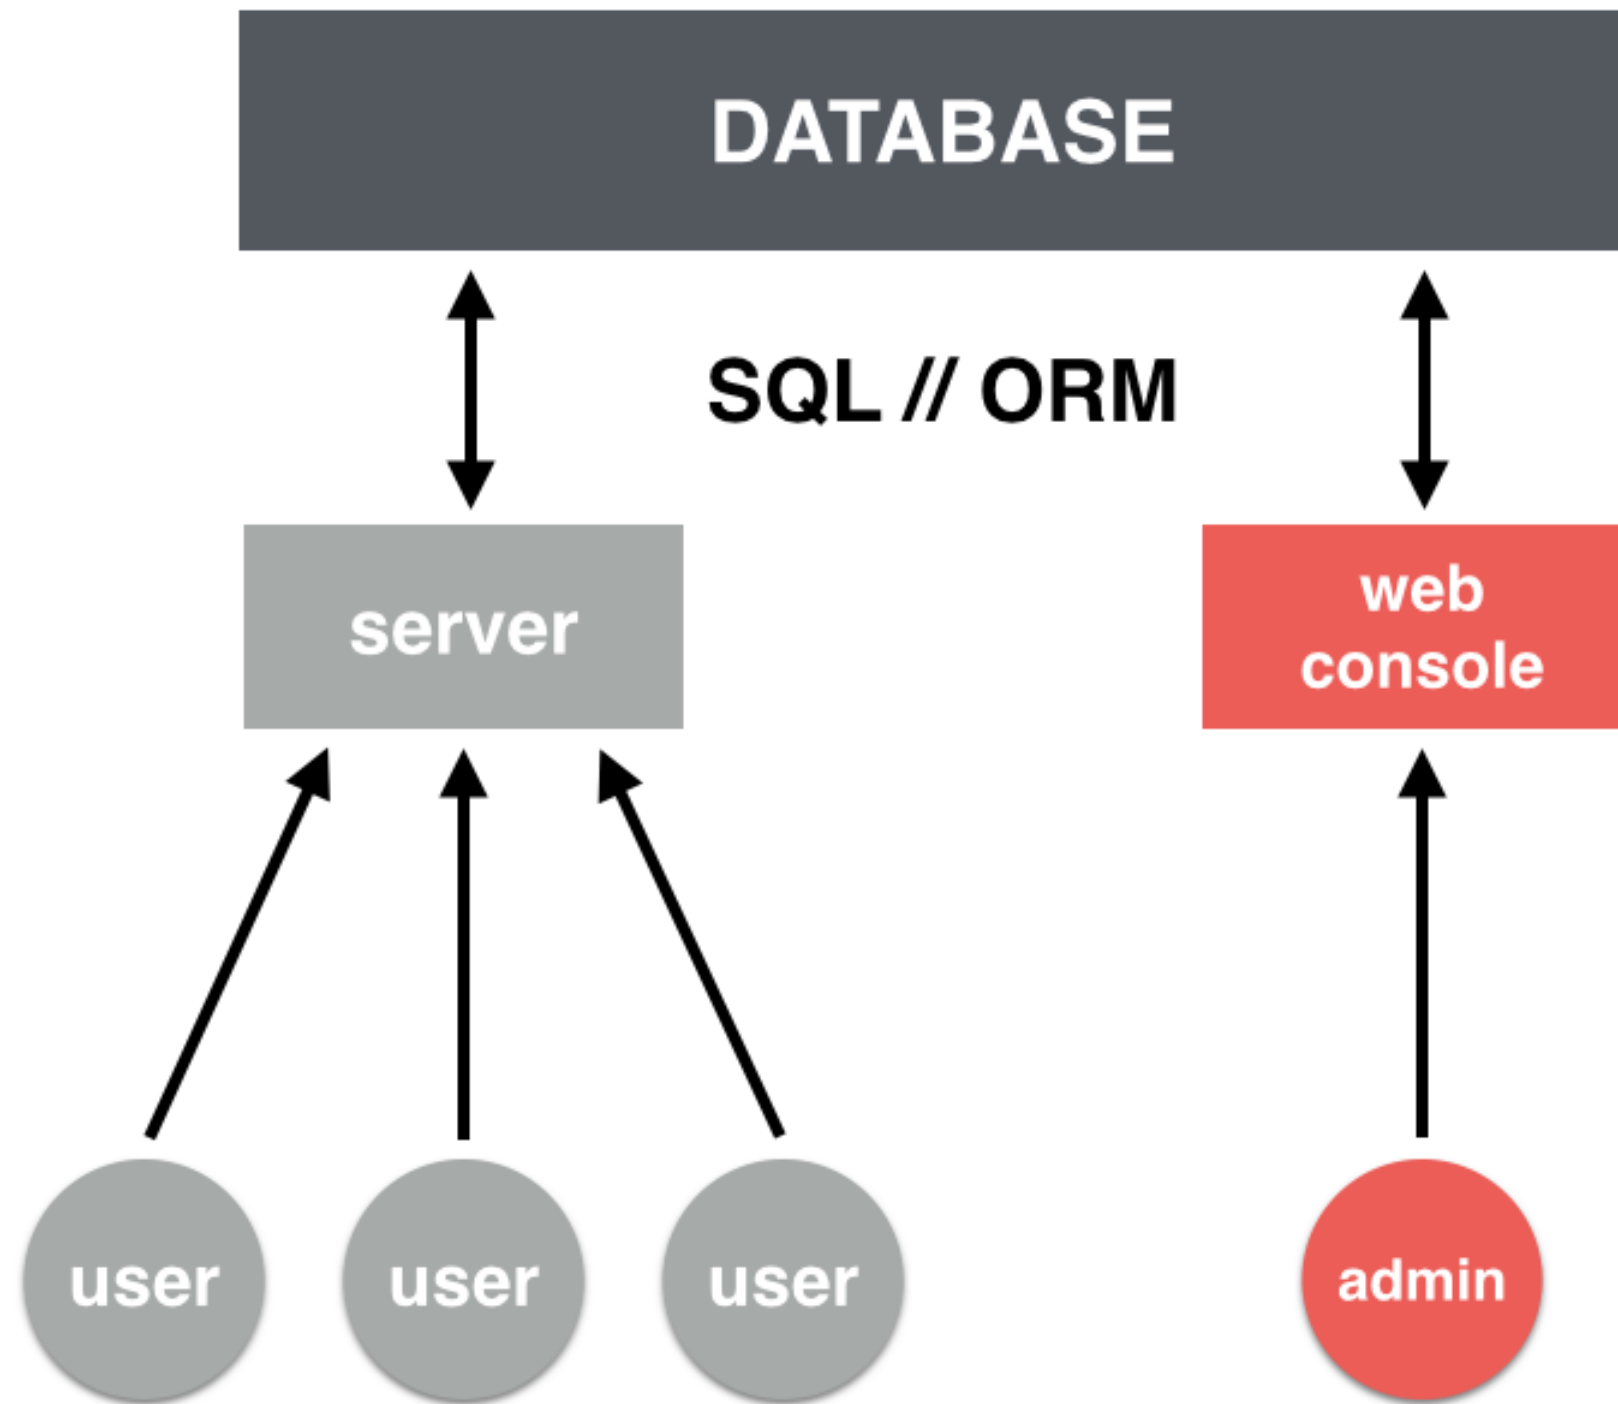
`twistd` is a daemon that helps run applications.


`twistd web --wsgi=dashboard.app`

Be careful when running as a service under upstart/systemd/init!

DEMO

**WHAT ABOUT DATA
PERSISTENCE?**





1. `twisted.enterprise.adbapi`

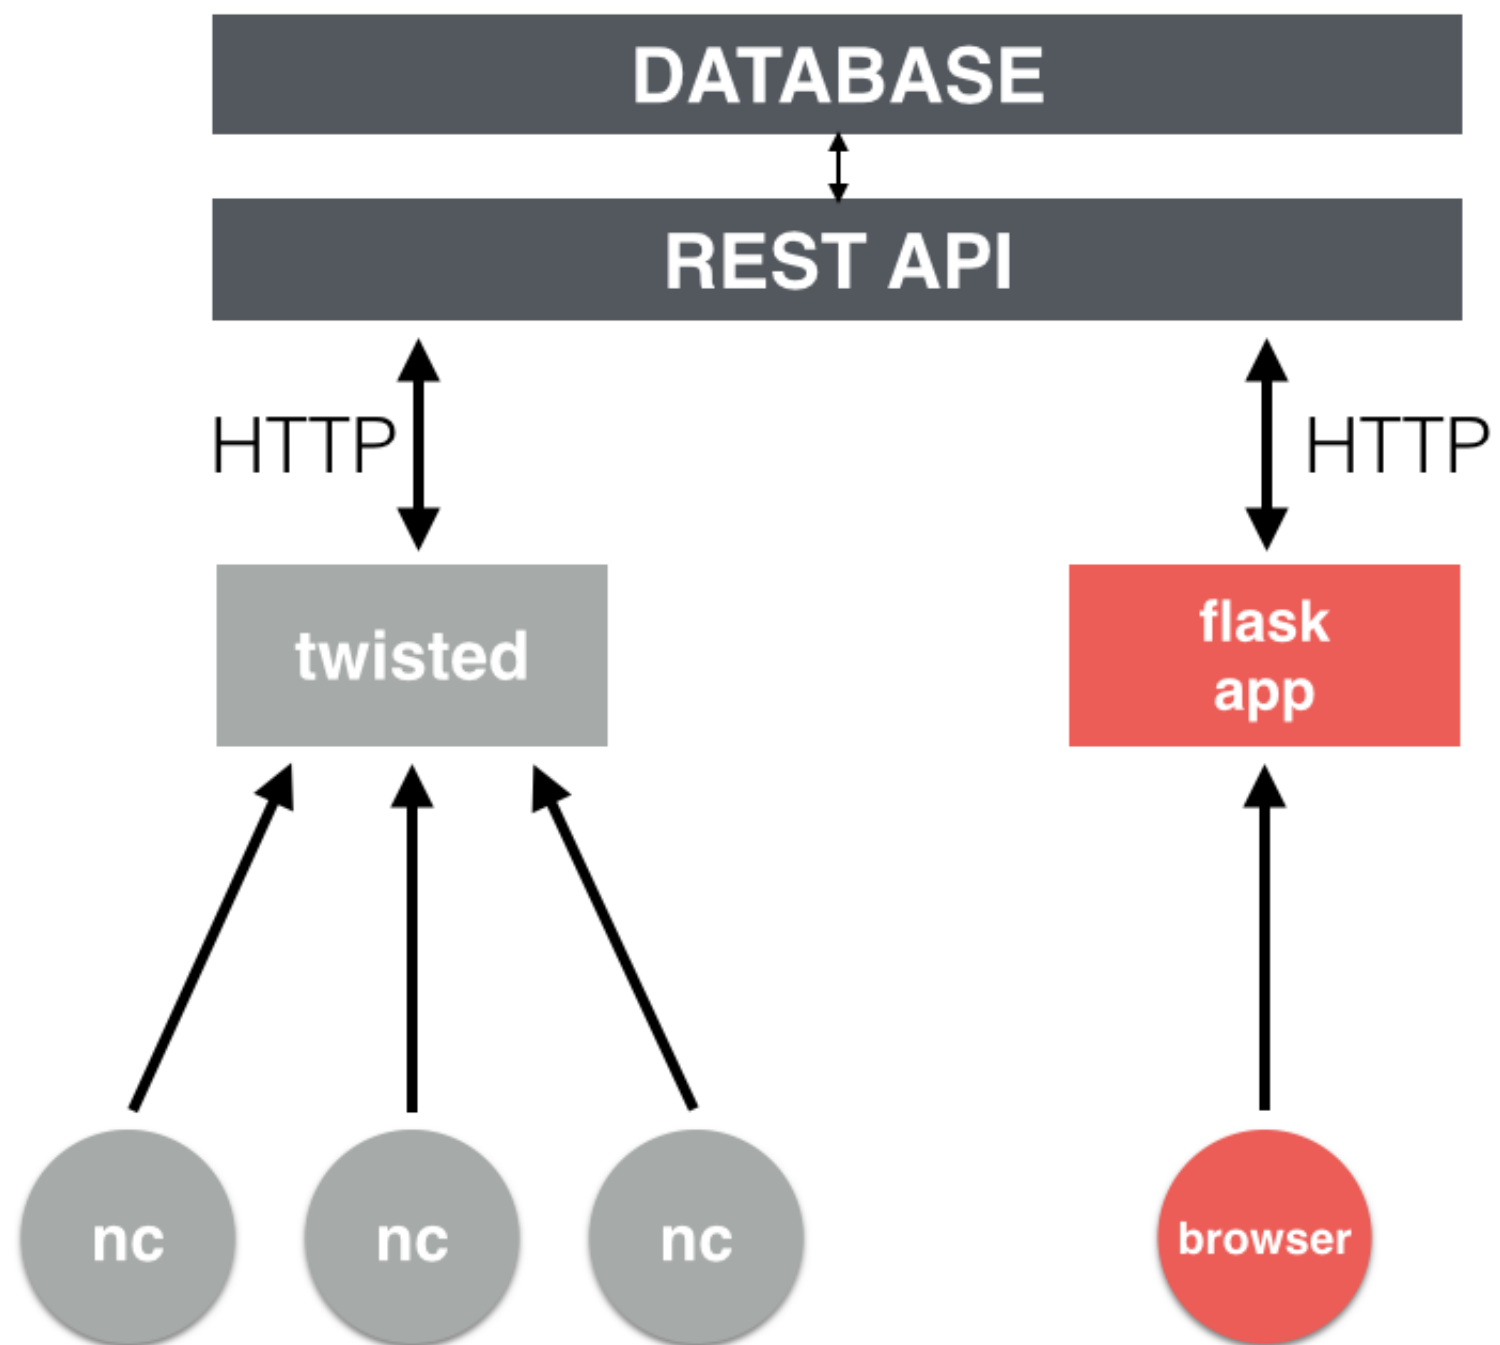
2. **SQLAlchemy**

3. **Django ORM**

...

CAREFUL NOT TO BLOCK!

**IS THERE A BETTER, MORE
REUSABLE WAY?**



**HOW WOULD YOU SCALE TO MANY TWISTED
SERVERS?**

- ▶ **focus on keeping state in one place**
- ▶ **make Twisted services stateless**
 - ▶ **load balance accordingly**

RECAP

- ▶ **what is async programming**
 - ▶ **what is Twisted**
- ▶ **when/why to use Twisted**

TWISTED RECAP

- ▶ **event loop (reactor)**
 - ▶ **deferreds**
 - ▶ **protocols (factories, transports)**
- ▶ **using Twisted to build larger systems/services**

WHAT TO RESEARCH NEXT

- ▶ **trial: testing, the twisted way**
- ▶ **inline callbacks: synchronous-looking deferreds in twisted**
 - ▶ **twistd: running twisted code as a service**
- ▶ **read http://krondo.com/?page_id=1327**

THANK YOU - QUESTIONS?

GITHUB.COM/PERCIPIENT/TALKS

patrick@percipientnetworks.com

stephen@percipientnetworks.com

(btw, we are sponsoring DjangoCon! Find us!)